Per population, in time order: training set | validation set (the last 93 days before the training cut) | 186-day gap | test set (from the test start to the last test origin + 31 days). Released recipe: training from scratch with early stopping on the validation set, then the validation span is returned to the training set for 8 more epochs (labels ending within 24 h of the cut excluded); nothing at or after the training cut is used for training.

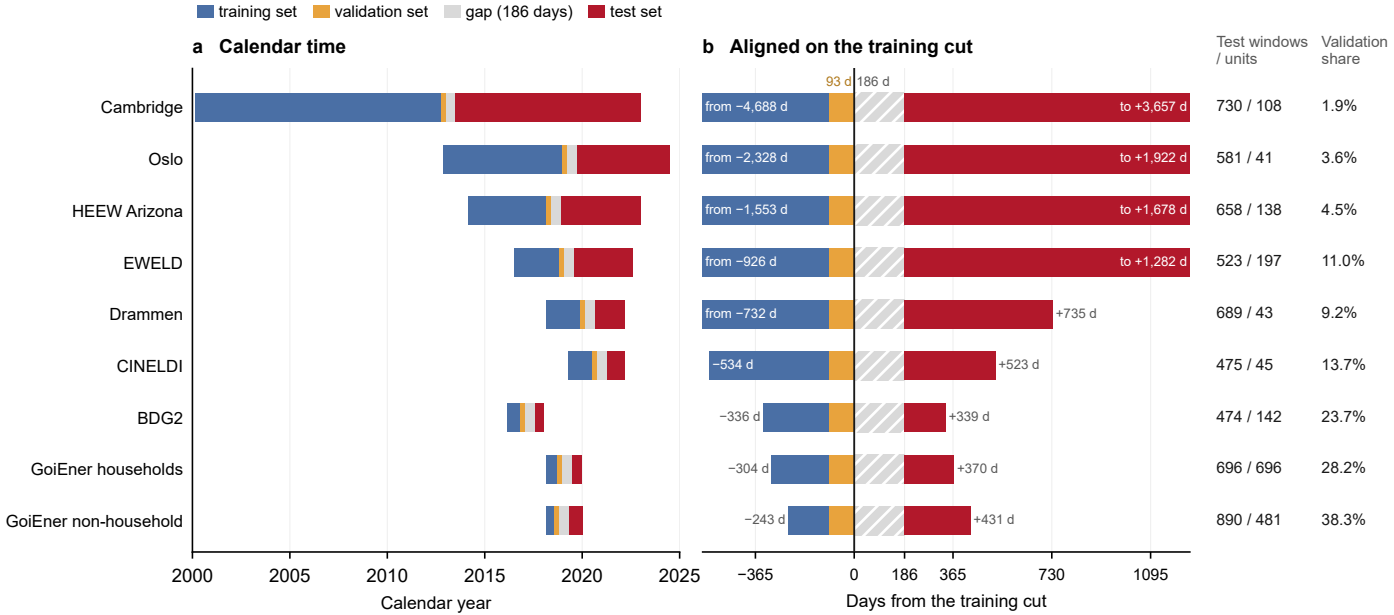

The training set also borrows windows of other populations that end before the validation start (not drawn). Validation share: the share of the population's own windows held out for early stopping.